

Home Assignment:

Build a Document Question Answering (Document QA) Agent

Develop an AI-powered Document QA agent that enables users to upload a document from their local system and ask questions based on its content. The agent should extract the document content, understand the user's query, and provide accurate answers using only the uploaded document as the source of information. If the requested information is not available in the document, the agent should respond appropriately instead of generating unsupported answers.

Expected Outcome: Users should be able to upload a document, ask multiple questions, and receive context-aware answers derived only from the uploaded document.

list of nodes :

- 1.Manual Trigger
- 2.Read /write files from disk
- 3.Extract from file
- 4.Basic LLM chain

Note: I have used DeepLearning.pdf that is provide by testleaf

My workflow 4 - n8n[DEV]

http://localhost:5678/workflow/Vt4kNq9ZqohdS5ow

Personal / My workflow 4

Editor Executions Evaluations

Publish

When clicking "Execute workflow"

Read/Write Files from Disk

Extract from File

Basic LLM Chain

Groq Chat Model

Execute workflow

Logs

Clear execution

Groq Chat Model Success in 2.04s 2,535 Tokens

Input

Output

Success in 2.942s 2,535 Tokens

When clicking "Execute work..."

Read/Write Files from Disk

Extract from File

Basic LLM Chain

Groq Chat Model

System: You are a Document Question Answering assistant. Answer the user's question using ONLY the document content provided below. Rules: 1. Use only information from the document. 2. Do not use your general knowledge. 3. Do not make assumptions. 4. Do not invent or hallucinate information. 5. If the answer is not available in the document, respond exactly: "The requested information is not available in the uploaded document." 6. Give a clear and concise answer. DOCUMENT CONTENT: Deep Learning & Generative AI: A comprehensive guide to understanding deep learning, neural networks, generative AI, and large language models — the foundational technologies reshaping how machines understand and create content. Deep Learning Generative AI LLMs CPU vs GPU What is Deep Learning? Deep Learning is a specialized subset of machine learning that uses neural networks with multiple layers — known as deep neural networks — to analyze and learn from vast amounts of data. Unlike traditional ML, deep learning models can automatically discover representations and features from raw data without explicit hand-engineering. Input Data: Raw data fed into the network (images, text, audio, etc.). Hidden Layers: Multiple layers of neurons that progressively extract higher-level patterns and features. Output: The final prediction, classification, or generated result produced by the model.

http://localhost:5678/workflow/Vt4kNq9ZqohdS5ow

Personal / My workflow 4

Editor Executions Evaluations

Publish

When clicking "Execute workflow"

Read/Write Files from Disk

Extract from File

Basic LLM Chain

Groq Chat Model

Execute workflow

Logs

Clear execution

Groq Chat Model Success in 2.04s 2,535 Tokens

Input

Output

Success in 2.942s 2,535 Tokens

When clicking "Execute work..."

Read/Write Files from Disk

Extract from File

Basic LLM Chain

Groq Chat Model

System: You are a Document Question Answering assistant. Answer the user's question using ONLY the document content provided below. Rules: 1. Use only information from the document. 2. Do not use your general knowledge. 3. Do not make assumptions. 4. Do not invent or hallucinate information. 5. If the answer is not available in the document, respond exactly: "The requested information is not available in the uploaded document." 6. Give a clear and concise answer. DOCUMENT CONTENT: Deep Learning & Generative AI: A comprehensive guide to understanding deep learning, neural networks, generative AI, and large language models — the foundational technologies reshaping how machines understand and create content. Deep Learning Generative AI LLMs CPU vs GPU What is Deep Learning? Deep Learning is a specialized subset of machine learning that uses neural networks with multiple layers — known as deep neural networks — to analyze and learn from vast amounts of data. Unlike traditional ML, deep learning models can automatically discover representations and features from raw data without explicit hand-engineering. 1 Input Data: Raw data fed into the network — images, text, audio, or structured records. 2 Hidden Layers: Multiple layers of neurons progressively extract higher-level patterns and features. 3 Output: The final prediction, classification, or generated result produced by the model. The AI Landscape: Where Deep Learning Fits: Deep learning sits within a broader hierarchy of AI disciplines. Understanding how these layers relate helps clarify what each approach is capable of and when it's the right tool to use. LLMs represent the cutting edge of this hierarchy — massive generative models built on deep neural architectures, capable of understanding and producing human language at scale. What is Generative AI? Generative AI refers to a class of artificial intelligence systems that can create entirely new content — including images, text, video, audio, and more — by learning patterns from existing datasets. Rather than simply classifying or predicting, generative models synthesize novel outputs that were never in the training data. 3. What is Generative AI? Generative AI is a class of artificial intelligence systems that can create entirely new content—such as images, text, video, audio, and more—by learning patterns from existing datasets. Instead of merely classifying or predicting, generative models synthesize novel outputs that were never present in the training data. 4. What is the difference between CPU and GPU? CPU (Central Processing Unit): 4-8 cores optimized for sequential, low-latency tasks. Excels at complex logic, branching, and interactive responsiveness. Best for general-purpose computing and serial processing. GPU (Graphics Processing Unit): Hundreds of thousands of cores designed for massively parallel workloads. Executes thousands of matrix operations simultaneously, providing high throughput. Essential for deep-learning training and other parallel-intensive tasks.

My workflow 4 - n8n[DEV]

http://localhost:5678/workflow/Vt4kNq9ZqohdS5ow

PersonalMy workflow 4

EditorExecutionsEvaluations

Execute workflow

Workflow Diagram

Logs

INPUT

OUTPUT

Overview

AI Assistant

Templates

Insights

Help

Settings

Success in 2.942s2,535 Tokens

When clicking "Execute work..."

Read/Write Files from Disk

Extract from File

Basic LLM Chain

Groq Chat Model

Success in 2.04s2,535 Tokens

System: You are a Document Question Answering assistant. Answer the user's question using ONLY the document content provided below. Rules: 1. Use only information from the document. 2. Do not use your general knowledge. 3. Do not make assumptions. 4. Do not invent or hallucinate information. 5. If the answer is not available in the document, respond exactly: "The requested information is not available in the uploaded document." 6. Give a clear and concise answer. DOCUMENT CONTENT: Deep Learning & Generative AI: A comprehensive guide to understanding deep learning, neural networks, generative AI, and large language models — the foundational technologies reshaping how machines understand and create content. Deep Learning Generative AI: LLMs CPU vs GPU What is Deep Learning? Deep Learning is a specialized subset of machine learning that uses neural networks with multiple layers — known as deep neural networks — to analyze and learn from vast amounts of data. Unlike traditional ML, deep learning models can automatically discover representations and features from raw data without explicit hand-engineering. 1 Input Data Raw data fed into the network — images, text, audio, or structured records. 2 Hidden Layers Multiple layers of neurons progressively extract higher-level patterns and features. 3 Output The final prediction, classification, or generated result produced by the model. The AI Landscape: Where Deep Learning Fits Deep learning sits within a broader hierarchy of AI disciplines. Understanding how these layers relate helps clarify what each approach is capable of and when it's the right tool to use. LLMs represent the cutting edge of this hierarchy — massive generative models built on deep neural architectures, capable of understanding and producing human language at scale. What is Generative AI? Generative AI refers to a class of artificial intelligence systems that can create entirely new content — including images, text, video, audio, and more — by learning patterns from existing datasets. Rather than simply classifying or predicting, generative models synthesize novel outputs that were never in the training data. The

4. What is the difference between CPU and GPU?

CPU (Central Processing Unit):

4-8 cores optimized for sequential, low-latency tasks.

Excels at complex logic, branching, and interactive responsiveness.

Best for general-purpose computing and serial processing.

GPU (Graphics Processing Unit):

Hundreds to thousands of cores designed for massively parallel workloads.

Executes thousands of matrix operations simultaneously, providing high throughput.

Essential for deep-learning training and other parallel-intensive tasks.

5. What is the price of the GPU?

The requested information is not available in the uploaded document.